

rec02_solutions

September 9, 2020

1 CS 1656 – Introduction to Data Science (Fall 2020)

1.1 Instructor: Alexandros Labrinidis / Teaching Assistant: E. Karageorgos

1.2 ### Additional Credits: Xiaoting Li, T. Arabghalizi, P. Pham, A. Shein, Z. Agha

In this recitation you will be learning pandas dataframe basics and plotting in Python. In the end you will see an example for NLTK that is related to your first assignment.

Packages you will need are, * pandas * matplotlib

First step is to import the packages above. If import fails, it means that the package is not installed and you need to install the package using "pip install package_name".

```
In [1]: import matplotlib.pyplot as plt
import pandas as pd
import datetime
import io
import requests
```

For the sake of interactive display in Jupyter, we will enable matplotlib inline.

```
In [2]: %matplotlib inline
```

1.3 Dataframe Basics

DataFrame is a 2-dimensional labeled data structure with columns of potentially different types. You can think of it like a spreadsheet or SQL table. DataFrame accepts many different kinds of input: *Dict of 1D ndarrays, lists, dicts, or Series* 2-D numpy.ndarray *Structured or record ndarray* A Series *Another DataFrame

Along with the data, you can optionally pass index (row labels) and columns (column labels) arguments.

Now what is a Series? Series is a one-dimensional labeled array capable of holding any data type (integers, strings, floating point numbers, Python objects, etc.). You can think of it as a 1-dimensional dataframe. Series objects can also have index. ### Creating a Dataframe We will start off by creating a dataframe from Weather Underground Data retrieved from the url below.

```
In [3]: r = io.StringIO(requests.get('http://data.cs1656.org/KPIT_Aug17.csv', verify=False).content)
df = pd.read_csv(r,
sep=',', engine='python', parse_dates=['EST'])
```

C:\ProgramData\Anaconda3\lib\site-packages\urllib3\connectionpool.py:858: InsecureRequestWarning
InsecureRequestWarning)

To display the top 'n' rows of the dataframe, use the head() command below. The default is 5 rows.

In [4]: df.head()

```
Out[4]:
```

| | EST | Max TemperatureF | Mean TemperatureF | Min TemperatureF | \ |
|---|------------|------------------|-------------------|------------------|---|
| 0 | 2017-08-01 | 84 | 73 | 62 | |
| 1 | 2017-08-02 | 84 | 75 | 65 | |
| 2 | 2017-08-03 | 84 | 75 | 65 | |
| 3 | 2017-08-04 | 83 | 73 | 62 | |
| 4 | 2017-08-05 | 74 | 66 | 58 | |

| | Max Dew PointF | MeanDew PointF | Min DewpointF | Max Humidity | Mean Humidity | \ |
|---|----------------|----------------|---------------|--------------|---------------|---|
| 0 | 65 | 61 | 58 | 93 | 67 | |
| 1 | 66 | 64 | 60 | 93 | 70 | |
| 2 | 66 | 64 | 62 | 93 | 71 | |
| 3 | 68 | 65 | 58 | 90 | 73 | |
| 4 | 61 | 54 | 49 | 87 | 68 | |

| | Min Humidity | ... | Mean Sea Level Press. | \ |
|---|--------------|-----|-----------------------|---|
| 0 | 40 | ... | 30.12 | |
| 1 | 47 | ... | 30.08 | |
| 2 | 48 | ... | 30.07 | |
| 3 | 55 | ... | 29.93 | |
| 4 | 49 | ... | 30.04 | |

| | Min Sea Level Press. | Max VisibilityMiles | Mean VisibilityMiles | \ |
|---|----------------------|---------------------|----------------------|---|
| 0 | 30.07 | 10 | 10 | |
| 1 | 30.04 | 10 | 10 | |
| 2 | 30.01 | 10 | 10 | |
| 3 | 29.82 | 10 | 9 | |
| 4 | 29.94 | 10 | 10 | |

| | Min VisibilityMiles | Max Wind SpeedMPH | Mean Wind SpeedMPH | \ |
|---|---------------------|-------------------|--------------------|---|
| 0 | 10 | 12 | 3 | |
| 1 | 10 | 13 | 3 | |
| 2 | 10 | 14 | 5 | |
| 3 | 2 | 32 | 9 | |
| 4 | 10 | 15 | 10 | |

| | Min Wind SpeedMPH | PrecipitationIn | Events |
|---|-------------------|-----------------|--------|
| 0 | 16 | 0 | NaN |
| 1 | 16 | 0 | NaN |
| 2 | 17 | 0 | NaN |

```

3          39          0.41  Rain , Thunderstorm
4          22          T          NaN

```

```
[5 rows x 21 columns]
```

Now to find all the column names of the dataframe and their data types, type the following command.

```
In [5]: df.dtypes
```

```

Out[5]: EST          datetime64[ns]
Max TemperatureF    int64
Mean TemperatureF   int64
Min TemperatureF    int64
Max Dew PointF      int64
MeanDew PointF      int64
Min DewpointF       int64
Max Humidity         int64
Mean Humidity        int64
Min Humidity         int64
Max Sea Level Press. float64
Mean Sea Level Press. float64
Min Sea Level Press. float64
Max VisibilityMiles  int64
Mean VisibilityMiles int64
Min VisibilityMiles  int64
Max Wind SpeedMPH    int64
Mean Wind SpeedMPH   int64
Min Wind SpeedMPH    int64
PrecipitationIn      object
Events               object
dtype: object

```

Notice the type of 'EST' column. We will find out why that's relevant a few steps later. ###
 Accessing Dataframe Columns There are two ways to access a dataframe column. The first way is accessing it like a dictionary as shown below. We will be using head function to show the first few rows only.

```
In [6]: df['EST'].head()
```

```

Out[6]: 0    2017-08-01
1    2017-08-02
2    2017-08-03
3    2017-08-04
4    2017-08-05
Name: EST, dtype: datetime64[ns]

```

The second way is to access using dot. But this only works if the column name is a valid variable name without any spacing.

```
In [7]: df.EST.head()
```

```
Out[7]: 0    2017-08-01
        1    2017-08-02
        2    2017-08-03
        3    2017-08-04
        4    2017-08-05
        Name: EST, dtype: datetime64[ns]
```

You can also access multiple columns by passing list of column names.

```
In [8]: df[['EST', 'Mean TemperatureF']].head()
```

```
Out[8]:
```

| | EST | Mean TemperatureF |
|---|------------|-------------------|
| 0 | 2017-08-01 | 73 |
| 1 | 2017-08-02 | 75 |
| 2 | 2017-08-03 | 75 |
| 3 | 2017-08-04 | 73 |
| 4 | 2017-08-05 | 66 |


1.4 Plotting

1.4.1 Basic Plot

Now lets start with basic plotting in Python first. We will use the plot function. Note that plot returns a tuple of handle and labels. If you need the plot handle in the future you will assign a variable to the plot function's return value.

```
In [9]: p1 = plt.plot(df['EST'],df['MeanDew PointF'])
        p2 = plt.plot(df['EST'],df['Mean TemperatureF'])
        plt.legend([p1[0],p2[0]], ['Mean Dew Point', 'Mean Temperature'])

        plt.show()
```



That does not look too pretty. Let's format the graph and plot again.


```
In [10]: # Initializing a larger figure
fig = plt.figure(figsize=(10, 6))

# Plotting
p1 = plt.plot(df['EST'],df['MeanDew PointF'])
p2 = plt.plot(df['EST'],df['Mean TemperatureF'])
plt.legend([p1[0],p2[0]], ['Mean Dew Point', 'Mean Temperature'])

# Formatting graph
plt.xticks(rotation = 90, fontsize = 8)
plt.xlabel('Date')
plt.ylabel('Mean Temperature')
plt.title('Mean Tempertaures for August 2017')

# Are we ready to show the formatted graph now? Not yet. Because we want
# to save our graph figure this time.In order to use the save command, it
# is important to save before the show command because the show command
# clears the axis of the figure after displaying.

plt.savefig("basic_plot.png")
plt.show()
```




1.4.2 Bar Plot

```
In [11]: fig = plt.figure(figsize=(10, 6))
plt.bar(range(len(df['EST'])),df['Mean Humidity'], align = 'center')

# Formatting graph
plt.xticks(range(len(df['EST'])), df['EST'].dt.strftime('%Y-%m-%d'),\
           rotation = 90, fontsize = 8)
plt.xlabel('Date')
plt.ylabel('Mean Humidity')
plt.title('Mean Humidity for August 2017')

plt.savefig("bar_plot.png")
plt.show()
```



Now, let's try to plot the two graphs above on the same figure using subplots. The code for plotting is the same as that shown above.

```
In [12]: fig = plt.figure(figsize=(10, 16))

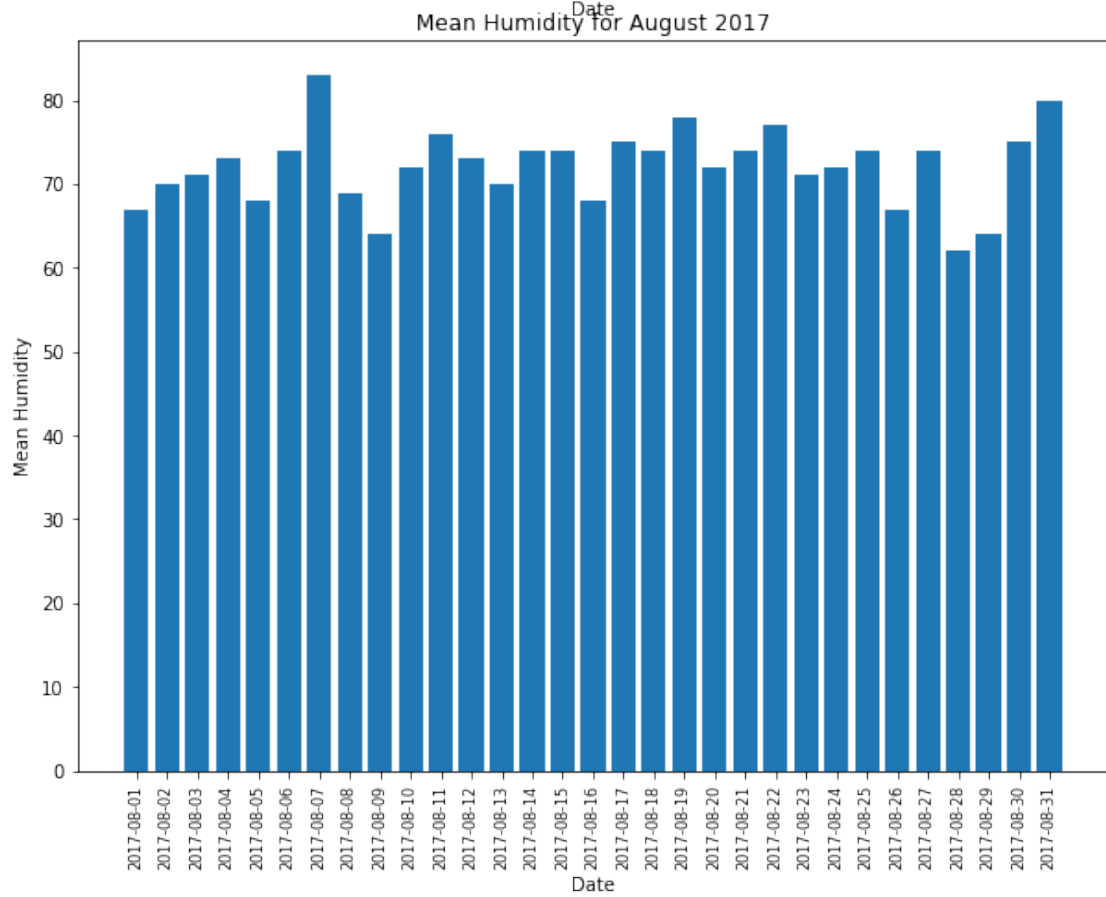

# Subplot of basic graph
plt.subplot(211)
p1 = plt.plot(df['EST'], df['Mean Dew PointF'])
p2 = plt.plot(df['EST'], df['Mean TemperatureF'])
plt.legend([p1[0], p2[0]], ['Mean Dew Point', 'Mean Temperature'])

plt.xticks(rotation = 90, fontsize = 8)
plt.xlabel('Date')
plt.ylabel('Mean Temperature')
plt.title('Mean Temperatures for August 2017')

# Subplot of bar graph
plt.subplot(212)
plt.bar(range(len(df['EST'])), df['Mean Humidity'], align = 'center')

plt.xticks(range(len(df['EST'])), df['EST'].dt.strftime('%Y-%m-%d'), \
            rotation = 90, fontsize = 8)
plt.xlabel('Date')
plt.ylabel('Mean Humidity')
```

```
plt.title('Mean Humidity for August 2017')  
  
plt.savefig("subplot_basic_bar.png")  
plt.show()
```

What is the advantage of subplots? It allows us to establish relationships between different variables and different statistics. Looking at the two subplots above, do you notice any relationship?

1.4.3 Scatter Plot


Scatter plots are commonly used to show correlation between variables. If the data points make a straight line going from the origin out to high x and high y values, then the variables are said to have a positive correlation. If the line goes from a high value on the y-axis down to a high value on the x-axis, the variables have a negative correlation. The closer the data points in the scatter plot, the higher the correlation between the two variables, or the stronger the relationship.

```
In [13]: fig = plt.figure(figsize=(10, 6))

plt.scatter(df['Max TemperatureF'],df['Min TemperatureF'])

# Formatting graph
plt.xlabel('Max Temperture')
plt.ylabel('Min Temperature')
plt.title('Min and Max Temperature for August 2017')

plt.savefig("scatter_plot.png")
plt.show()
```



The graph above shows us that min and max temperature have a strong positive correlation.

2 NLTK and Stemming:

Stemming usually refers to a crude heuristic process that chops off the ends of words in the hope of achieving this goal correctly most of the time, and often includes the removal of derivational affixes.

```
In [14]: import nltk
         nltk.download('punkt')
         mysentence = """During a manned mission to Mars, Astronaut Mark Watney
                        (Matt Damon) is presumed dead after a fierce storm and left behind
                        by his crew.But Watney has survived and finds himself stranded and alone on

         mysentencetokens = nltk.word_tokenize(mysentence)
         #Stemming
         porter = nltk.PorterStemmer()
         loop = 0
         for token in mysentencetokens:
             mysentencetokens[loop] = porter.stem(token)
             loop += 1
         print ("Stemmed -->")
         print (mysentencetokens)

[nltk_data] Downloading package punkt to
[nltk_data] C:\Users\Angello\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!
Stemmed -->
['dure', 'a', 'man', 'mission', 'to', 'mar', ',', 'astronaut', 'mark', 'watney', '(', 'matt', 'd
```

2.1 Tasks

For your tasks, the input file is available at <http://data.cs1656.org/top12cities.csv>. The file consists of population density estimates and land area of several cities in USA.

You need to read the file into a dataframe and perform the following three tasks during the lab.

Task 1

Plot a scatter plot with 'land area' on the x-axis and '2014 estimate' on the y-axis. After observing the plot, do you think the two variables are strongly or weakly correlated? Is the correlation positive or negative?

Task 2

Plot a bar plot showing each city's 2014 population estimate given by '2014 estimate' column.

Task 3

Now that you plotted a simple bar plot, try plotting a grouped bar plot that shows both 2010 and 2014 estimate for each city on the same plot. This means that there will be two grouped bars per city on your graph.